



**David Eccles**  
**School of Business**  

---

**THE UNIVERSITY OF UTAH**

# **KNN Tutorial**

Data Mining for Business Intelligence

# OUTLINE

- KNN in R with IBk
- IBk package defaults and parameters.
- Exploring Train vs Test set results
- Cross-Validation Performance of various argument settings.
- Grid-Search and IBk best models
- Ablation Analysis

# IBk – Package Introduction

See IBk documentation and parameter settings at

<http://weka.sourceforge.net/doc.dev/weka/classifiers/lazy/IBk.html>

Parameters:

-I

Weight neighbours by the inverse of their distance (use when  $k > 1$ )

-F

Weight neighbours by  $1 - \text{their distance}$  (use when  $k > 1$ )

-K <number of neighbors>

Number of nearest neighbours (k) used in classification. (Default = 1)

-X

Select the number of nearest neighbours between 1 and the k value specified using hold-one-out evaluation on the training data (use when  $k > 1$ )

-E

Minimize mean squared error rather than mean absolute error when using -X option with numeric prediction.

# Data Import and Categorical Factorization

- Here we import the needed libraries
- Import the required dataset
- Create variables for use throughout the tutorial such as the metrics list and the seed value.
- Explore the structure and summarize the dataset.

```
44 library(caret)
45 library(Rweka)
46 library(rminer)
47 library(matrixStats)
48 library(knitr)
49 library(tictoc)
50 library(tidyverse)
51 tic()
52
53 # import data
54 cloud_wd <- getwd()
55 setwd(cloud_wd)
56 insurance <- read.csv(file = "insurance.csv", stringsAsFactors = TRUE)
57
58 ### Set up cv parameters
59
60 df <- insurance
61 target <- 7
62 seedval <- 500
63 metrics_list <- c("MAE", "RMSE", "MAPE", "RMSPE")
64 ```
65
66 ```{r}
67 str(insurance)
68 ```
69
70 ```{r}
71 summary(insurance)
72 ```
```

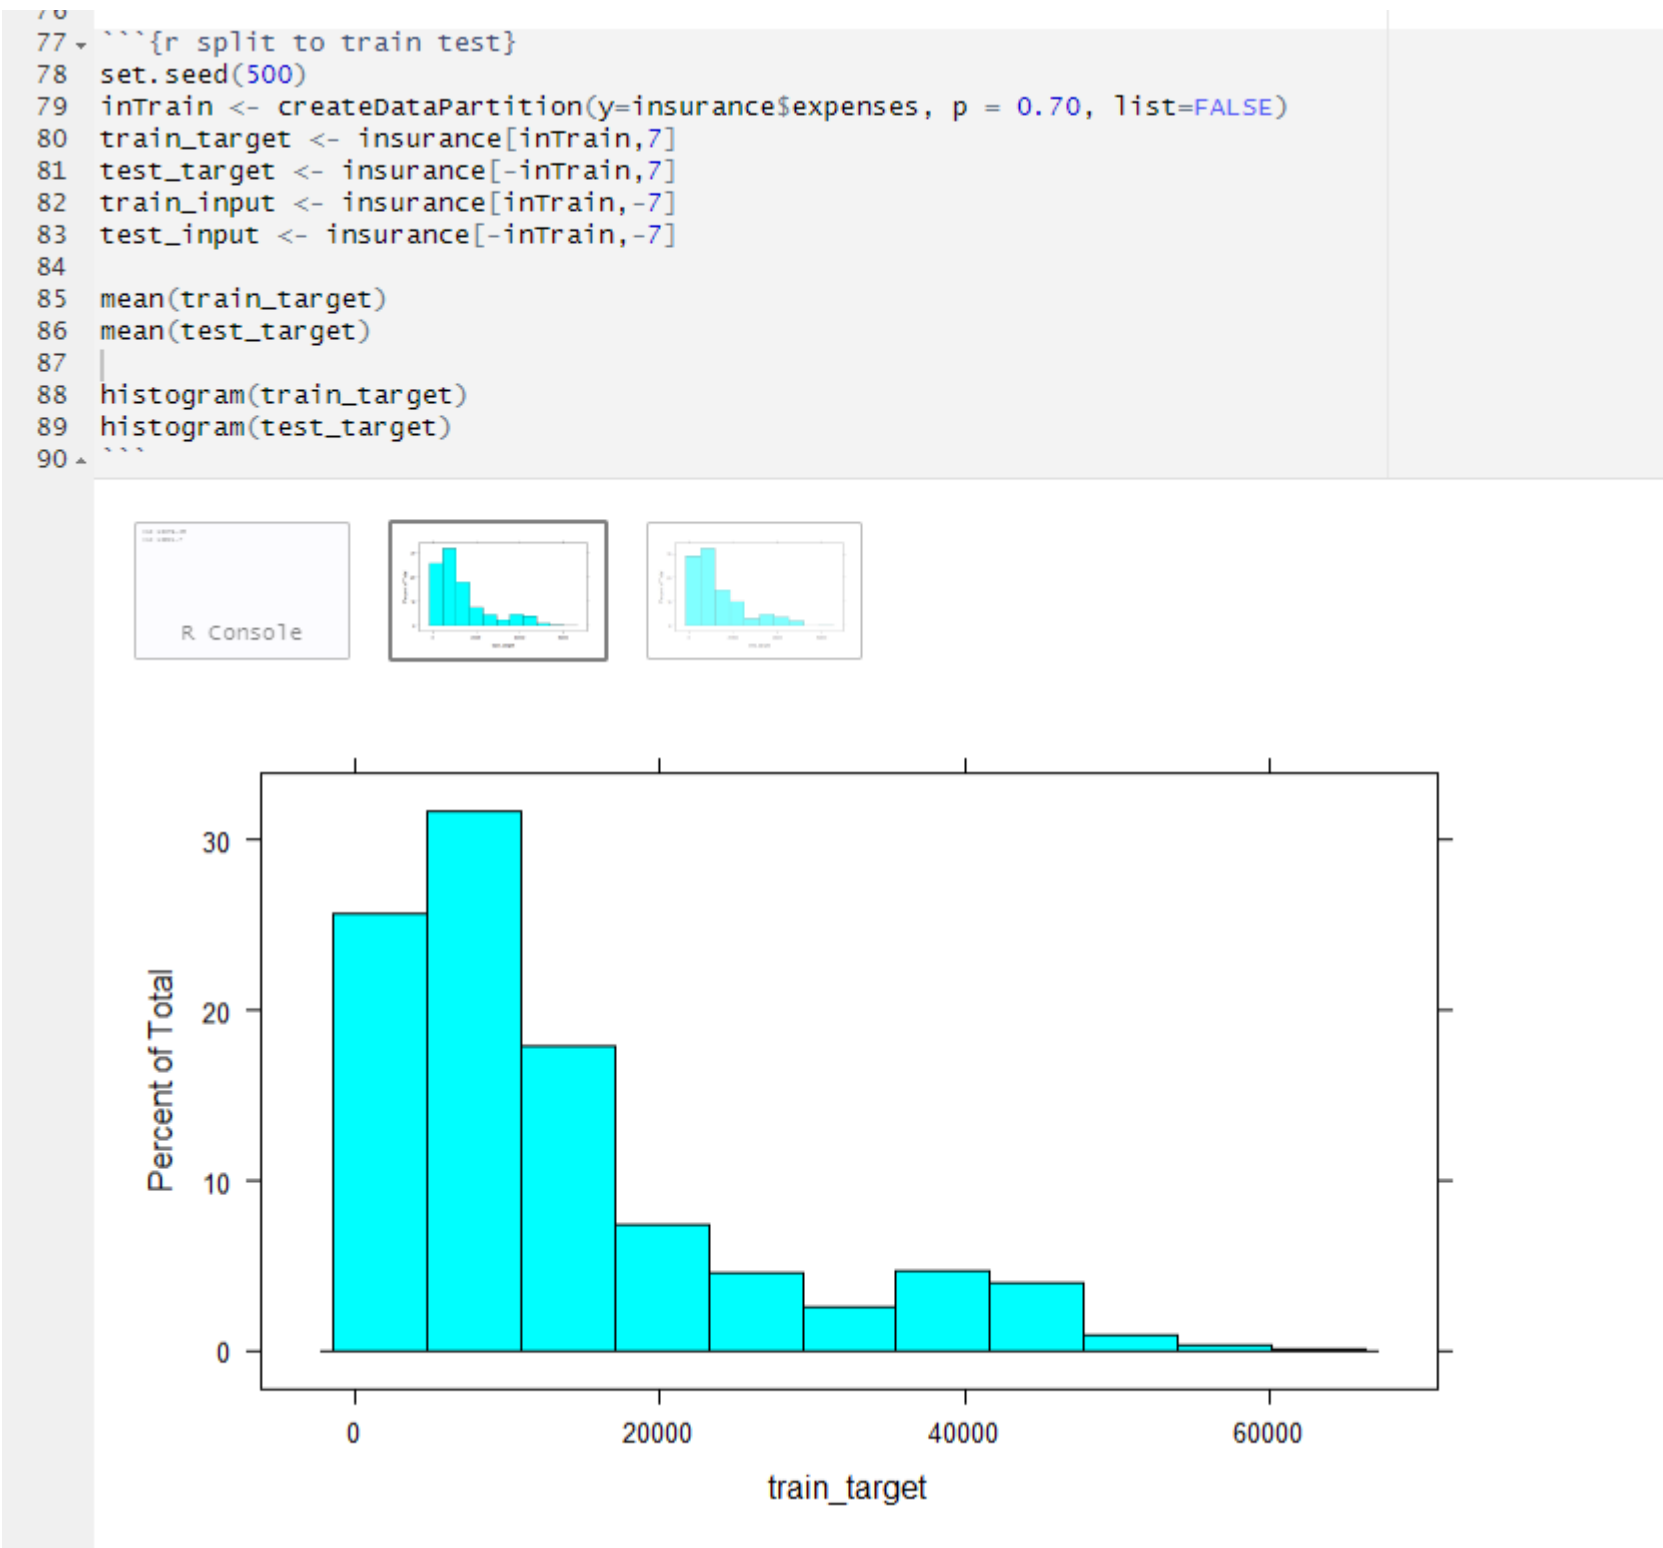
'data.frame': 1338 obs. of 7 variables:

```
$ age      : int  19 18 28 33 32 31 46 37 37 60 ...
$ sex      : Factor w/ 2 levels "female","male": 1 2 2 2 2 1 1 1 2 1 ...
$ bmi      : num  27.9 33.8 33 22.7 28.9 25.7 33.4 27.7 29.8 25.8 ...
$ children: int   0 1 3 0 0 0 1 3 2 0 ...
$ smoker   : Factor w/ 2 levels "no","yes": 2 1 1 1 1 1 1 1 1 1 ...
$ region   : Factor w/ 4 levels "northeast","northwest",...: 4 3 3 2 2 3 3 2 1 2 ...
$ expenses: num  16885 1726 4449 21984 3867 ...
```

age	sex	bmi	children	smoker	region	expenses
Min. :18.00	female:662	Min. :16.00	Min. :0.000	no :1064	northeast:324	Min. : 1122
1st Qu.:27.00	male :676	1st Qu.:26.30	1st Qu.:0.000	yes: 274	northwest:325	1st Qu.: 4740
Median :39.00		Median :30.40	Median :1.000		southeast:364	Median : 9382
Mean :39.21		Mean :30.67	Mean :1.095		southwest:325	Mean :13270
3rd Qu.:51.00		3rd Qu.:34.70	3rd Qu.:2.000			3rd Qu.:16640
Max. :64.00		Max. :53.10	Max. :5.000			Max. :63770

# Data Partition – Simple Holdout

- Split the data into test and train sets, extract the target variable into a new variable for both sets.
- Show the distribution of the target in both train and test to prove both sets are similar when comparing via the target.



# K = 1 on Train and Test

- Compare model performance from Train to Test using a K = 1 (1 nearest neighbor setting).
- Train set errors are 0 because each instance can be predicted using itself resulting in no errors at all.
- Test set errors are much higher. The nearest neighbor in Test is not the exact same observation.

```
96 when using K = 1 on the training set we will then achieve an MAE = 0 (or if classification 100% accuracy)
97
98 ## MAE for K = 1
99
100 ```{r Simple IBk examples}
101
102 # K = 1
103
104 knn_model_k_1 <- IBk(train_target ~ ., data = train_input, control = weka_control(k=1))
105 knn_model_k_1
106 ```
```

IB1 instance-based classifier using 1 nearest neighbour(s) for classification			
MAE	RMSE	MAPE	RMSPE
0	0	0	0
MAE	RMSE	MAPE	RMSPE
3803.667750	7217.775045	37.651870	9.425661

# K = 5 on Train and Test

- Higher errors on both Train (expected) and Test.
- It appears K = 5 using unweighted voting is not as good as K = 1.

```
120
121 ## MAE for K = 5
122
123 ```{r}
124 knn_model_k_5 <- IBk(train_target ~ ., data = train_input, control = weka_control(k=5))
125 knn_model_k_5
126 ```

IB1 instance-based classifier
using 5 nearest neighbour(s) for classification

127
128 ```{r k5 train}
129 insample_pred <- predict(knn_model_k_5, train_input)
130 mmetric(train_target, insample_pred, metrics_list)
131 ```

      MAE      RMSE      MAPE      RMSPE
2827.042013 4698.609321  26.490765   5.015548

132
133 ```{r k5 test}
134 test_pred <- predict(knn_model_k_5, test_input)
135 mmetric(test_target, test_pred, metrics_list)
136 ```

      MAE      RMSE      MAPE      RMSPE
3994.90551 6388.61020  41.18298   7.17385

137
```



# K = 1 and I = TRUE on Train and Test

- The errors shouldn't change from using only k=1 since we are using only 1 neighbor. Weighting is useful when  $k > 1$ .

```
139 ## weighted distance
140
141 ## MAE for K = 1 and I = TRUE
142
143 ```{r Simple IBk examples I=TRUE}
144
145 # K = 1
146
147 knn_model_k1_ittrue <- IBk(train_target ~ ., data = train_input, control = weka_control(K=1, I=TRUE))
148 knn_model_k1_ittrue
149 ```
```

IB1 instance-based classifier using 1 inverse-distance-weighted nearest neighbour(s) for classification			
<hr/>			
150			
151 ```{r k1 train I=TRUE}			
152 insample_pred <- predict(knn_model_k1_ittrue, train_input)			
153 mmetric(train_target, insample_pred, metrics_list)			
154 ```			
<hr/>			
MAE RMSE MAPE RMSPE			
0 0 0 0			
<hr/>			
155			
156 ```{r k1 test I=TRUE}			
157 test_pred <- predict(knn_model_k1_ittrue, test_input)			
158 mmetric(test_target, test_pred, metrics_list)			
159 ```			
<hr/>			
MAE RMSE MAPE RMSPE			
3803.667750 7217.775045 37.651870 9.425661			
<hr/>			



# K = 5 and I = TRUE on Train and Test

- Weighted voting has improved Train performance due to more weight on the exact same instance being the most heavily weighted neighbor.
- In test we've also improved our performance as compared to K = 5 without weighted voting (I=FALSE)

```
163
164 ▾ ## MAE for K = 5 and I = TRUE
165
166 ▾ ```{r}
167 knn_model_k5_ittrue <- IBk(train_target ~ ., data = train_input, control = weka_control(K=5, I=TRUE))
168 knn_model_k5_ittrue
169 ▾ ```
```

IB1 instance-based classifier using 5 inverse-distance-weighted nearest neighbour(s) for classification			
--	--	--	--

```
170
171 ▾ ```{r k5 train I=TRUE}
172 insample_pred <- predict(knn_model_k5_ittrue, train_input)
173 mmetric(train_target, insample_pred, metrics_list)
174 ▾ ```
```

MAE	RMSE	MAPE	RMSPE
239.6097554	475.9304220	2.6135382	0.7020244

```
175
176 ▾ ```{r k5 test I=TRUE}
177 test_pred <- predict(knn_model_k5_ittrue, test_input)
178 mmetric(test_target, test_pred, metrics_list)
179 ▾ ```
```

MAE	RMSE	MAPE	RMSPE
3821.653457	6260.476947	39.918157	7.330998

# K = from 1 to 40 and I = TRUE and X=TRUE on Train and Test

- We now ask the model to eval k values from 1 to 40 using leave one out cross validation. The best K will be selected.
- The model has chosen a K = 3 based on its ability to reduce the MAE.

```
184 ## MAE for K = 40, I = TRUE, X = TRUE.
185
186 Try k from 1 to 40, using inverse of distance weights.
187
188 Demonstrating how to use the Leave one out cross-validation built into the IBk package to determine the best k.
189
190 ```{r X TRUE}
191 knn_model_itrue_xtrue <- IBk(train_target ~ ., data = train_input, control = weka_control(K=40, I=TRUE, X=TRUE))
192 knn_model_itrue_xtrue
193 ```
```

IB1 instance-based classifier using 3 inverse-distance-weighted nearest neighbour(s) for classification			
--	--	--	--

```
194 The model has chosen k = 3 to minimize the MAE.
195
196 ```{r}
197 insample_pred <- predict(knn_model_itrue_xtrue, train_input)
198 mmetric(train_target, insample_pred, metrics_list)
199 ```
```

MAE	RMSE	MAPE	RMSPE
159.9757828	372.7041571	1.8370163	0.6670268

```
200
201 ```{r}
202 test_pred <- predict(knn_model_itrue_xtrue, test_input)
203 mmetric(test_target, test_pred, metrics_list)
204 ```
```

MAE	RMSE	MAPE	RMSPE
3813.619044	6464.275066	40.250375	8.281032

# IBk custom cross-validation function

- Create a custom cross validation function which allows us to explore k and I values.
- In addition, we change the end of this function to include columns for the k and i parameters as well as return a dataframe instead of a kable.

```
216  
217 ~~~{r Define cv_IBk}  
218  
219 cv_IBk <- function(df, target, nFolds, seedval, metrics_list, k, i)  
220 {  
221   # create folds using the assigned values  
222  
223   set.seed(seedval)  
224   folds = createFolds(df[,target],nFolds)  
225  
226   # The lapply loop  
227  
228   cv_results <- lapply(folds, function(x)  
229   {  
230     # data preparation:  
231  
232     test_target <- df[x,target]  
233     test_input <- df[x,-target]  
234  
235     train_target <- df[-x,target]  
236     train_input <- df[-x,-target]  
237     pred_model <- IBk(train_target ~ .,data = train_input,control = weka_control(K=k,I=i))  
238     pred <- predict(pred_model, test_input)  
239     train_pred <- predict(pred_model, train_input)  
240  
241     return(mmetric(test_target,pred,metrics_list))  
242   })  
243  
244   cv_results_m <- as.matrix(as.data.frame(cv_results))  
245   cv_mean<- as.matrix(rowMeans(cv_results_m))  
246   cv_sd <- as.matrix(rowSds(cv_results_m))  
247   colnames(cv_mean) <- "Mean"  
248   colnames(cv_sd) <- "sd"  
249   cv_df <- data.frame(t(cbind(cv_mean,cv_sd))) %>% round(2)  
250   cv_df$param_K <- k  
251   cv_df$param_I <- as.logical(i)  
252   cv_df <- cv_df %>% rownames_to_column(var = "measure")  
253   cv_df  
254 }  
255 ~~~  
256
```

# IBk custom cross-validation function

- We run the new `cv_function` with various values of `k` and `I = FALSE`

```
262 ~~~{r call cv_IBk_train and cv_IBk}
263 cv_IBk(df, target, 3, seedVal, metrics_list, 1, FALSE)|
264 ~~~
```

Description: df [2 x 7]						
measure <chr>	MAE <dbl>	RMSE <dbl>	MAPE <dbl>	RMSPE <dbl>	param_K <dbl>	param_I <lgl>
Mean	3728.35	7169.88	39.04	11.91	1	FALSE
Sd	218.95	302.90	3.17	1.54	1	FALSE
2 rows						

```
265
266
267 ## Call cross-validation with k=2 and I=FALSE
268 ~~~{r}
269 cv_IBk(df, target, 3, seedVal, metrics_list, 2, FALSE)
270 ~~~
```

Description: df [2 x 7]						
measure <chr>	MAE <dbl>	RMSE <dbl>	MAPE <dbl>	RMSPE <dbl>	param_K <dbl>	param_I <lgl>
Mean	3497.67	6212.26	38.02	9.81	2	FALSE
Sd	68.79	88.36	3.54	1.55	2	FALSE
2 rows						

```
271
272 ## Call cross-validation with k=3 and I=FALSE
273 ~~~{r}
274 cv_IBk(df, target, 3, seedVal, metrics_list, 3, FALSE)
275 ~~~
```

Description: df [2 x 7]						
measure <chr>	MAE <dbl>	RMSE <dbl>	MAPE <dbl>	RMSPE <dbl>	param_K <dbl>	param_I <lgl>
Mean	3575.87	5986.29	38.02	8.01	3	FALSE
Sd	115.05	202.05	3.05	0.96	3	FALSE
2 rows						

# Parameter Grid Setups

- Build a grid of k values from 1 to 40 and l options (TRUE,FALSE)
- Notice our dataframe is 80 rows long.

```
313
314 # Part 3.
315
316 Use a custom grid search to find our best hyperparameter combination.
317
318 ## Build the grid
319
320 ```{r}
321 # Create multiple vectors
322 param_k <- c(seq(1,40))
323 param_l <- c(FALSE, TRUE)
324
325 # Generate a grid of all combinations
326 grid <- expand.grid(param_k, param_l, stringsAsFactors = FALSE)
327
328 colnames(grid) <- c("k","l")
329 # Print the grid
330
331 grid
332 ^
```


Description: df [80 x 2]

k <int>	l <lgl>
1	FALSE
2	FALSE
3	FALSE
4	FALSE
5	FALSE
6	FALSE
7	FALSE
8	FALSE
9	FALSE
10	FALSE

1-10 of 80 rows

# Parameter Grid Looping function

- Define a new named function called “run\_grid\_cv” which loops over each row in the dataframe and calls the cross validation function for the rows setups for I and K.
- This function will loop 80 times based on the previous grid setups.



```
334 ## define a named function to loop through the grid
335
336 This function will take the grid as input and call the cross validation function for each row in the grid. Finally, each cross validation
337 deviation will be returned as a dataframe.
338 ```{r}
339
340 run_grid_cv <- function(grid) {
341   results <- data.frame()
342   for (i in 1:nrow(grid)) {
343     row <- grid[i,] # Get the i-th row
344     cv_result <- cv_IBk(df = df, target = target, nFolds = 5, seedVal = seedVal, metrics_list = metrics_list, k = row$k, i = row$i)
345     results <- rbind(results, cv_result)
346   }
347   return(results)
348 }
349 ```
```

# Calling the grid search function

- We simply pass the grid as a parameter into the new named function.
- The function takes the grid and loops over each row calling the cross validation function each time.
- We store the resulting dataframe in a variable named `cv_grid_results`.

```
351 ## run the grid search, return a dataframe with the results.  
352  
353 ```{r}  
354 cv_grid_results <- run_grid_cv(grid)  
355 ```  
356
```



# Filter the result dataframe. Find the best 10 MAE.

- We now demonstrate how to filter the dataframe to find only the Mean results, and arrange the dataframe in ascending order to show the lowest MAE. We use the head function to show only the best 10 MAE results. k = 2 is the winner.

```
358 {r}
359 # pull the top 3 based on MAE
360
361 cv_grid_results %>%
362   filter(measure == 'Mean') %>% |
363   arrange(MAE) %>%
364   head(10)
365 }
```

Description: df [10 x 7]

	measure <chr>	MAE <dbl>	RMSE <dbl>	MAPE <dbl>	RMSPE <dbl>	param_K <int>	param_I <lgl>
1	Mean	3261.69	6012.29	32.22	8.18	2	TRUE
2	Mean	3325.64	5850.02	33.11	7.40	3	TRUE
3	Mean	3357.56	6021.81	33.29	8.00	2	FALSE
4	Mean	3366.06	5684.82	33.17	6.46	5	TRUE
5	Mean	3371.33	5772.18	33.48	6.92	4	TRUE
6	Mean	3388.12	5681.46	33.24	6.21	6	TRUE
7	Mean	3416.70	5697.26	33.10	6.00	7	TRUE
8	Mean	3456.47	5719.96	33.15	5.85	8	TRUE
9	Mean	3474.73	5743.77	33.10	5.72	9	TRUE
10	Mean	3483.40	5925.34	34.63	7.18	3	FALSE

1-10 of 10 rows

# Filter the result dataframe. Find the best 10 MAPE.

- We now demonstrate how to filter the dataframe to find only the Mean results, and arrange the dataframe in ascending order to show the lowest MAPE. We use the head function to show only the best 10 MAPE results. k = 2 is the winner.
- It's interesting to note that some of the top models are not the same based on MAE vs MAPE.

```
367 {r}
368 # pull the top 10 based on MAPE
369
370 cv_grid_results %>%
371   filter(measure == 'Mean') %>%
372   arrange(MAPE) %>%
373   head(10)
374 }
```

Description: df [10 × 7]

	measure <chr>	MAE <dbl>	RMSE <dbl>	MAPE <dbl>	RMSPE <dbl>	param_K <int>	param_I <lgl>
1	Mean	3261.69	6012.29	32.22	8.18	2	TRUE
2	Mean	3484.70	5757.88	32.80	5.57	10	TRUE
3	Mean	3498.20	5764.89	32.82	5.49	11	TRUE
4	Mean	3516.06	5795.59	33.04	5.45	12	TRUE
5	Mean	3416.70	5697.26	33.10	6.00	7	TRUE
6	Mean	3474.73	5743.77	33.10	5.72	9	TRUE
7	Mean	3325.64	5850.02	33.11	7.40	3	TRUE
8	Mean	3456.47	5719.96	33.15	5.85	8	TRUE
9	Mean	3366.06	5684.82	33.17	6.46	5	TRUE
10	Mean	3388.12	5681.46	33.24	6.21	6	TRUE

1-10 of 10 rows

# Ablation Analysis

- Do we see performance drop or improve when we remove a single column from the predictors?
- Here we remove each column one at a time and use cross validation to examine the results.
- We run a baseline first, then remove age, then sex, then bmi, then children then smoker and finally region.
- Based on MAE removing Region appears to lower the MAE, however it increases the MAPE.

385  
386  
387  
388  
389  
390  
391  
392  
393  
394  
395  
396  
397  
398  
399  
400  
401  
402  
403

# Part 4

## Ablation Analysis (test feature removal)

Here we systematically remove each feature from the model to see the impact. This allows us to get a sense of feature importance. Some features cause the errors to rise substantially. Removing region actually improves the model.

```
```{r}
cv_IBk(df, 7, 5, seedval, metrics_list, 2, TRUE) # baseline
cv_IBk(df[, -c(1)], 6, 5, seedval, metrics_list, 2, TRUE)
cv_IBk(df[, -c(2)], 6, 5, seedval, metrics_list, 2, TRUE)
cv_IBk(df[, -c(3)], 6, 5, seedval, metrics_list, 2, TRUE)
cv_IBk(df[, -c(4)], 6, 5, seedval, metrics_list, 2, TRUE)
cv_IBk(df[, -c(5)], 6, 5, seedval, metrics_list, 2, TRUE)
cv_IBk(df[, -c(6)], 6, 5, seedval, metrics_list, 2, TRUE) # removing region appears to improve the model. slightly.
```
```

data.frame  
2 x 7

data.frame  
2 x 7

data.frame  
2 x 7

data.frame  
2 x 7

data.frame  
2 x 7

data.frame  
2 x 7

data.frame  
2 x 7

Description: df [2 x 7]

| measure<br><chr> | MAE<br><dbl> | RMSE<br><dbl> | MAPE<br><dbl> | RMSPE<br><dbl> | param_K<br><dbl> | param_I<br><lg> |
|------------------|--------------|---------------|---------------|----------------|------------------|-----------------|
| Mean             | 3261.69      | 6012.29       | 32.22         | 8.18           | 2                | TRUE            |
| Sd               | 290.10       | 478.84        | 4.63          | 1.55           | 2                | TRUE            |

2 rows

394  
395  
396  
397  
398  
399  
400  
401  
402  
403

```{r}

cv\_IBk(df, 7, 5, seedval, metrics\_list, 2, TRUE) # baseline

cv\_IBk(df[, -c(1)], 6, 5, seedval, metrics\_list, 2, TRUE)

cv\_IBk(df[, -c(2)], 6, 5, seedval, metrics\_list, 2, TRUE)

cv\_IBk(df[, -c(3)], 6, 5, seedval, metrics\_list, 2, TRUE)

cv\_IBk(df[, -c(4)], 6, 5, seedval, metrics\_list, 2, TRUE)

cv\_IBk(df[, -c(5)], 6, 5, seedval, metrics\_list, 2, TRUE)

cv\_IBk(df[, -c(6)], 6, 5, seedval, metrics\_list, 2, TRUE) # removing region appears to improve the model. slightly.

```

data.frame  
2 x 7

data.frame  
2 x 7

data.frame  
2 x 7

data.frame  
2 x 7

data.frame  
2 x 7

data.frame  
2 x 7

data.frame  
2 x 7

Description: df [2 x 7]

measure <chr>	MAE <dbl>	RMSE <dbl>	MAPE <dbl>	RMSPE <dbl>	param_K <dbl>	param_I <lg>
Mean	3192.73	6050.70	35.21	8.90	2	TRUE
Sd	241.30	377.82	5.22	1.34	2	TRUE

2 rows

Information Systems

Matt Pecsok



**David Eccles**  
**School of Business**

---

THE UNIVERSITY OF UTAH